

# Neural Networks

# Neural Networks

- Why neural networks?
- Model representation
- Examples and intuitions
- **Multi-class classification**

# Multiple output units: One-vs-all



Pedestrian



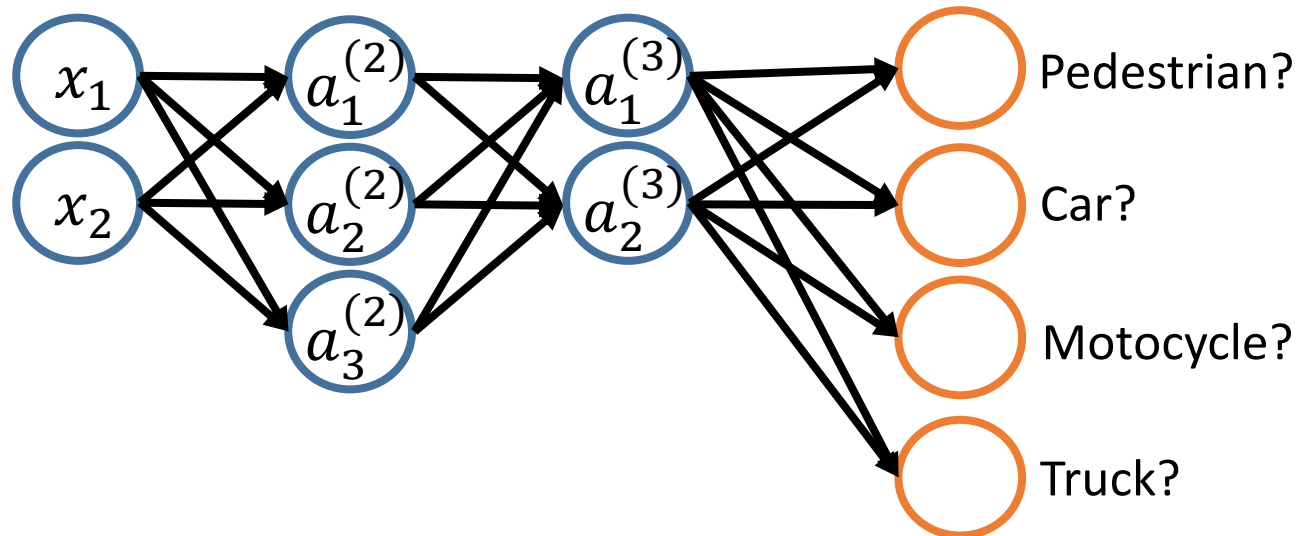
Car



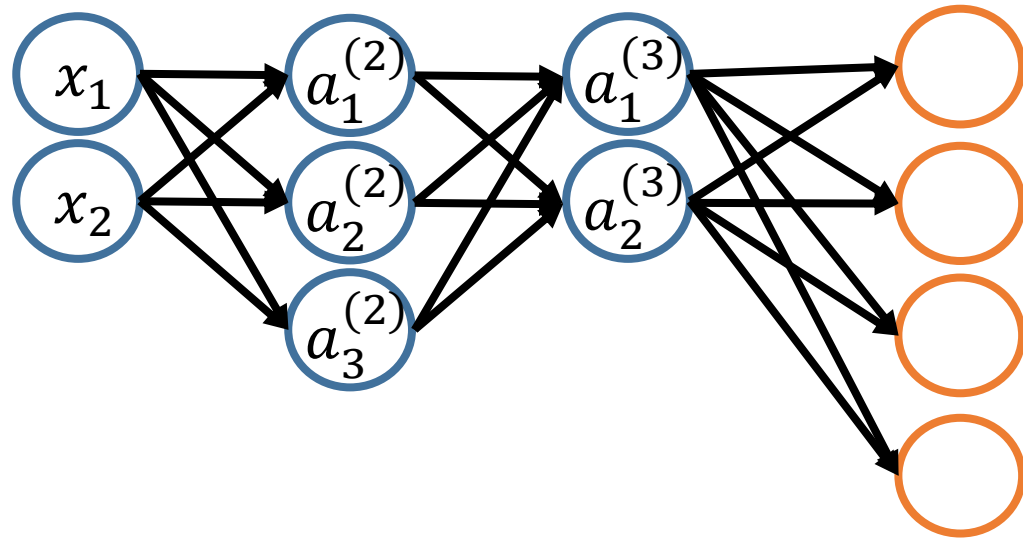
Motorcycle



Truck



# Multiple output units: One-vs-all

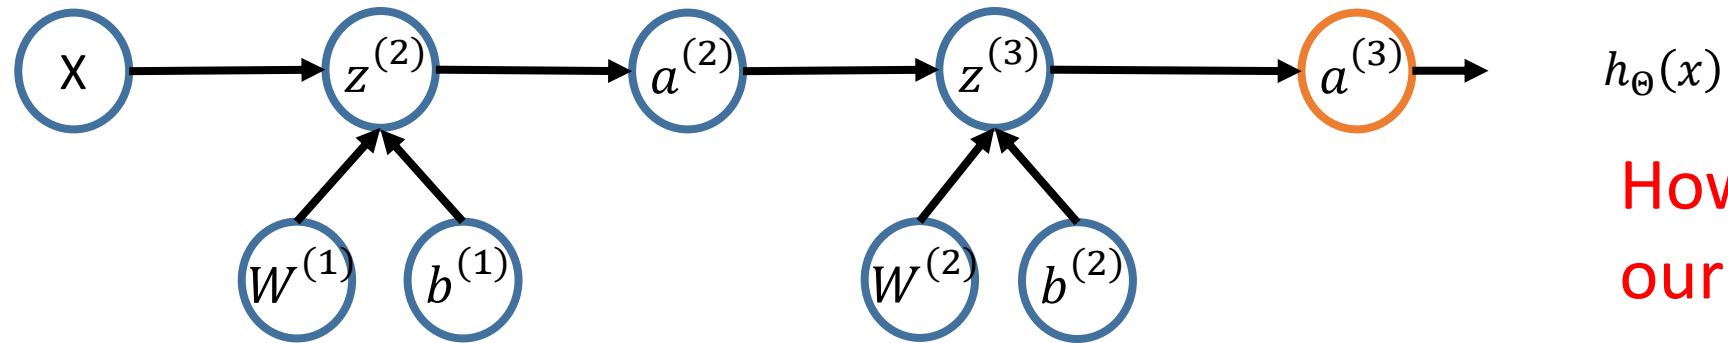


$$h_{\Theta}(x) \in R^4$$

Training set :  $(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots (x^{(m)}, y^{(m)})$ ,

$$y^{(i)} \text{ one of } \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

# Flow graph - Forward propagation



How do we evaluate  
our prediction?

$$z^{(2)} = \Theta^{(1)} x = \Theta^{(1)} a^{(1)}$$

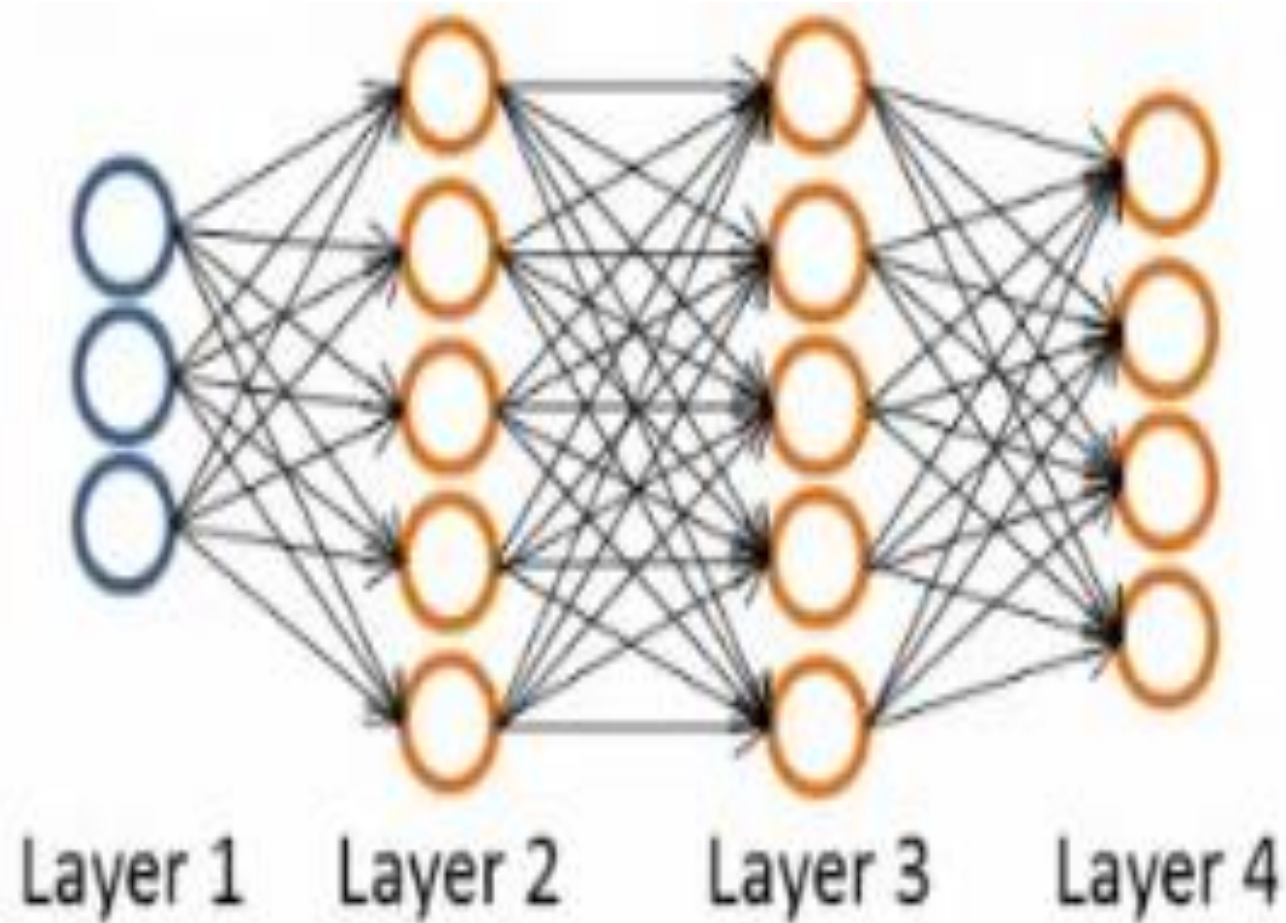
$$a^{(2)} = g(z^{(2)})$$

$$\text{Add } a_0^{(2)} = 1$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$h_{\Theta}(x) = a^{(3)} = g(z^{(3)})$$

# Neural Network cost function



# Cost function

Logistic regression:

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Neural network:

$$h_{\Theta}(x) \in \mathbb{R}^K \quad (h_{\Theta}(x))_i = i^{th} \text{ output}$$

$$J(\Theta) = -\frac{1}{m} \left[ \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h_{\Theta}(x^{(i)}))_k + (1 - y_k^{(i)}) \log(1 - (h_{\Theta}(x^{(i)}))_k) \right] \\ + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\theta_{ji}^{(l)})^2$$

# Cost function

- Our cost function now outputs a  $k$  dimensional vector
  - $h_{\theta}(x)$  is a  $k$  dimensional vector, so  $h_{\theta}(x)_i$  refers to the  $i$ th value in that vector
- Cost function  $J(\Theta)$  is
  - $[-1/m]$  times a sum of a similar term to which we had for logic regression
  - But now this is also a sum from  $k = 1$  through to  $K$  ( $K$  is number of output nodes)
    - Summation is a sum over the  $k$  output units - i.e. for each of the possible classes
    - So if we had 4 output units then the sum is  $k = 1$  to 4 of the logistic regression over each of the four output units in turn



# Gradient computation

$$J(\Theta) = -\frac{1}{m} \left[ \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(h_{\Theta}(x^{(i)}))_k + (1 - y_k^{(i)}) \log(1 - (h_{\Theta}(x^{(i)}))_k) \right] \\ + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{ji}^{(l)})^2$$

$$\min_{\Theta} J(\Theta)$$

Need to compute:

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$$

- This is a massive regularization summation term, it's a fairly straightforward triple nested summation.
- This is also called a weight decay term, as before, the lambda value determines the importance of the two halves
- The regularization term is similar to that in logistic regression. So, we have a cost function, but how do we minimize this

# Gradient computation

Given one training example  $(x, y)$

$$a^{(1)} = x$$

$$z^{(2)} = \Theta^{(1)} a^{(1)}$$

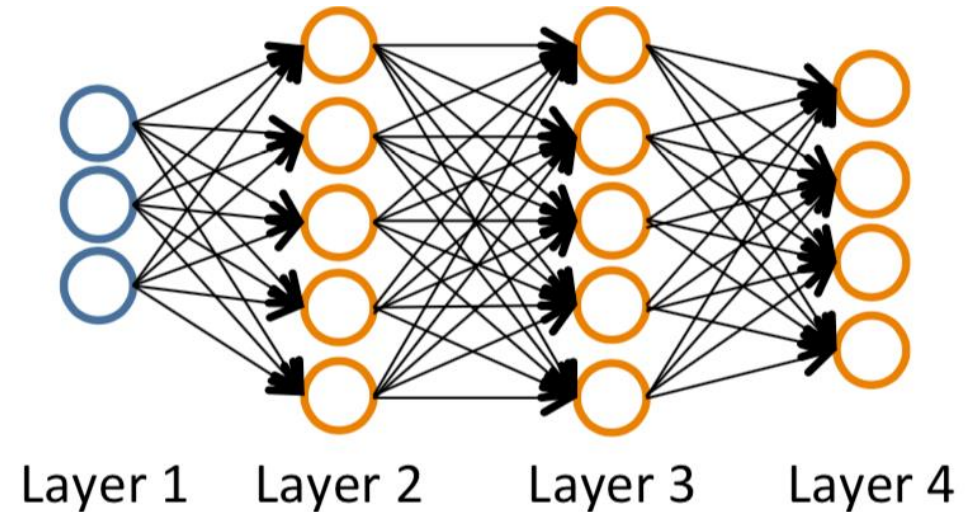
$$a^{(2)} = g(z^{(2)}) (\text{add } a_0^{(2)})$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$a^{(3)} = g(z^{(3)}) (\text{add } a_0^{(3)})$$

$$z^{(4)} = \Theta^{(3)} a^{(3)}$$

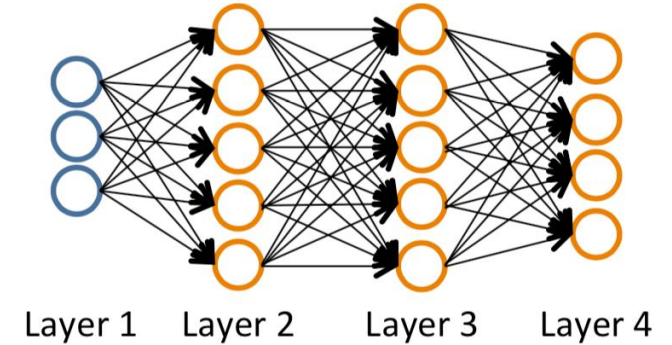
$$a^{(4)} = g(z^{(4)}) = h_{\Theta}(x)$$



# Gradient computation: Backpropagation

Intuition:  $\delta_j^{(l)}$  = “error” of node  $j$  in layer  $l$

For each output unit (layer  $L = 4$ )



$$\delta^{(4)} = a^{(4)} - y$$

$$\delta^{(3)} = \delta^{(4)} \frac{\partial \delta^{(4)}}{\partial z^{(3)}} = \delta^{(4)} \frac{\partial \delta^{(4)}}{\partial a^{(4)}} \frac{\partial a^{(4)}}{\partial z^{(4)}} \frac{\partial z^{(4)}}{\partial a^{(3)}} \frac{\partial a^{(3)}}{\partial z^{(3)}}$$

$$\delta^{(3)} = (\Theta^2)^\top \delta^3 \cdot *(a^{(3)} \cdot *(1 - a^{(3)}))$$

$$\delta^{(2)} = (\Theta^2)^\top \delta^2 \cdot *(a^{(2)} \cdot *(1 - a^{(2)}))$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$a^{(3)} = g(z^{(3)})$$

$$z^{(4)} = \Theta^{(3)} a^{(3)}$$

$$a^{(4)} = g(z^{(4)})$$

# Why do we do this?

- We do all this to get all the  $\delta$  terms, and we want the  $\delta$  terms because through a very complicated derivation you can use  $\delta$  to get the partial derivative of  $\Theta$  with respect to individual parameters (if you ignore regularization, or regularization is 0)

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = \mathbf{a}_j^l \delta_i^{(l+1)}$$

- By doing back propagation and computing the delta terms you can then compute the **partial derivative terms**
- We need the partial derivatives to minimize the cost function!

# Backpropagation algorithm

Training set  $\{(x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})\}$

Set  $\Theta^{(1)} = 0$

For  $i = 1$  to  $m$

Set  $a^{(1)} = x$

Perform forward propagation to compute  $a^{(l)}$  for  $l = 2..L$

use  $y^{(i)}$  to compute  $\delta^{(L)} = a^{(L)} - y^{(i)}$

Compute  $\delta^{(L-1)}, \delta^{(L-2)} \dots \delta^{(2)}$

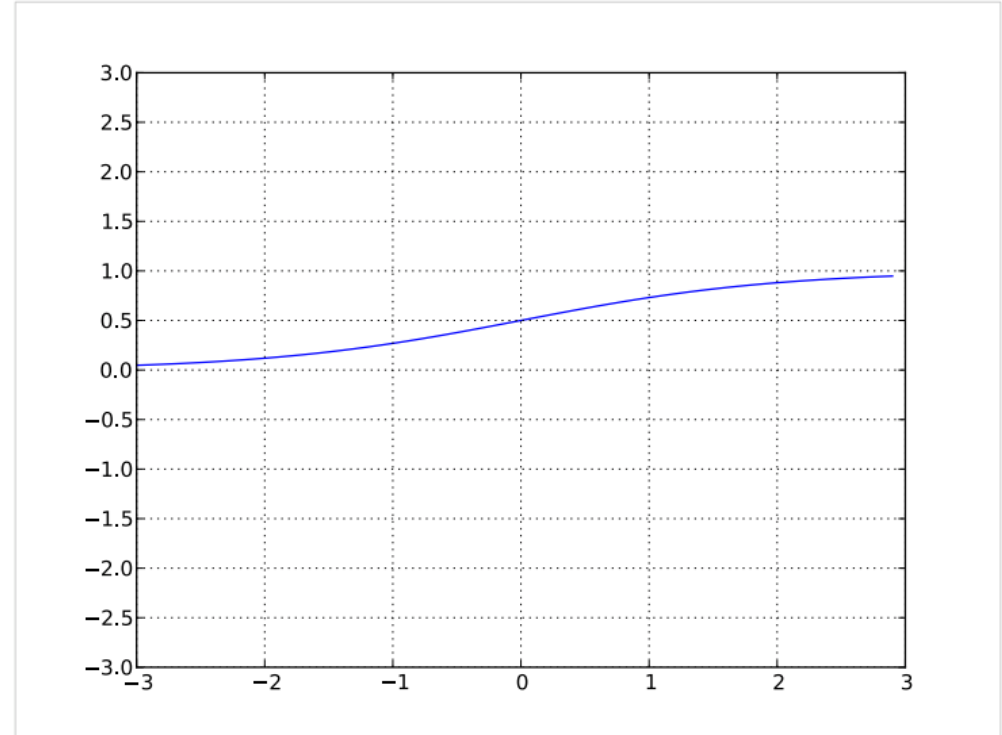
$$\Theta^{(l)} = \Theta^{(l)} - a^{(l)} \delta^{(l+1)}$$

# Activation - sigmoid

- Partial derivative

$$g'(x) = g(x)(1 - g(x))$$

- Squashes the neuron's pre-activation between 0 and 1
- Always positive
- Bounded
- Strictly increasing



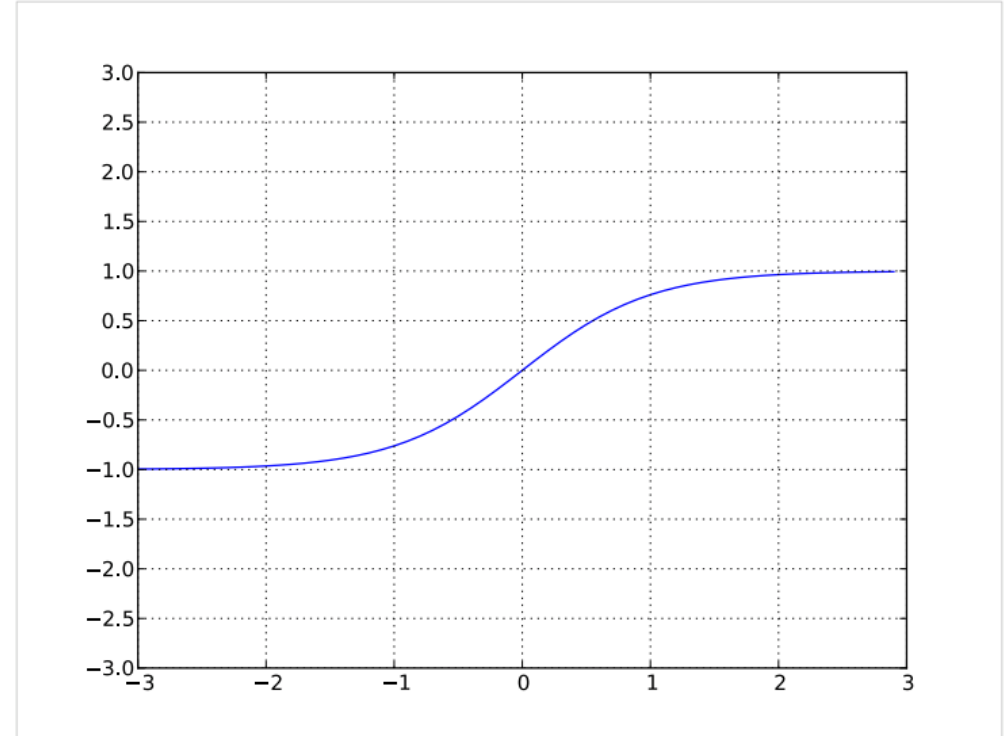
$$g(x) = \frac{1}{1 + e^{-x}}$$

# Activation - hyperbolic tangent (tanh)

- Partial derivative

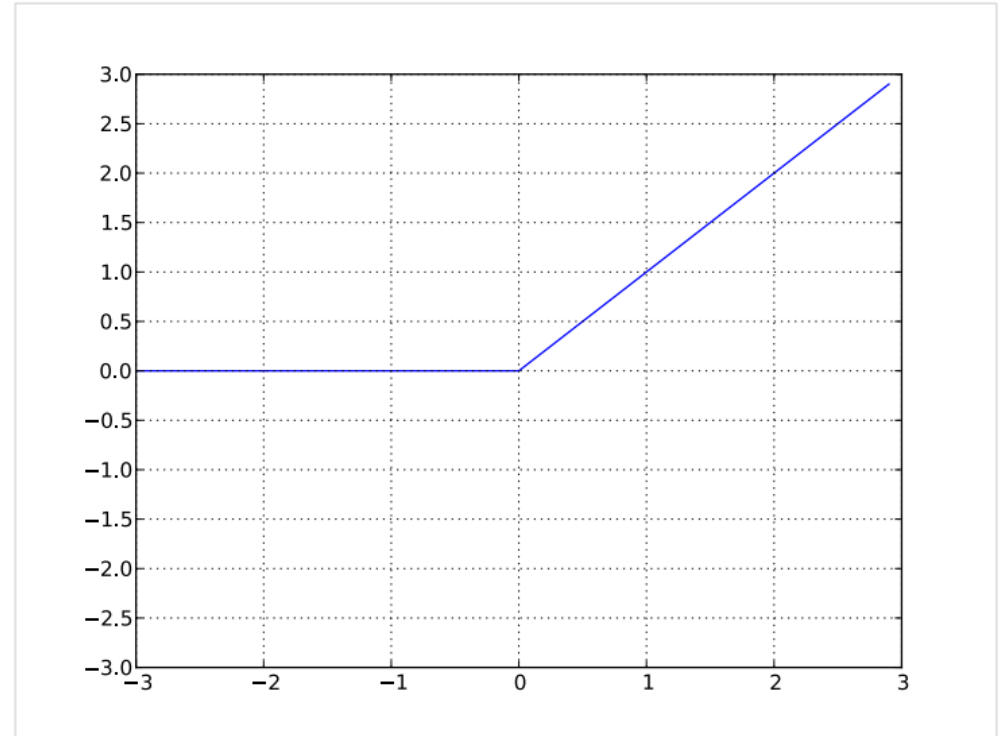
$$g'(x) = 1 - g(x)^2$$

- Squashes the neuron's pre-activation between -1 and 1
- Can be positive or negative
- Bounded
- Strictly increasing



$$g(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

# Activation - rectified linear(relu)



$$g(x) = \text{relu}(x) = \max(0, x)$$



# Backpropagation algorithm

Training set  $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$

Set  $\Delta_{ij}^{(l)} = 0$  (for all  $l, i, j$ ).

For  $i = 1$  to  $m$

Set  $a^{(1)} = x^{(i)}$

Perform forward propagation to compute  $a^{(l)}$  for  $l = 2, 3, \dots, L$

Using  $y^{(i)}$ , compute  $\delta^{(L)} = a^{(L)} - y^{(i)}$

Compute  $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$

$\Delta_{ij}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$

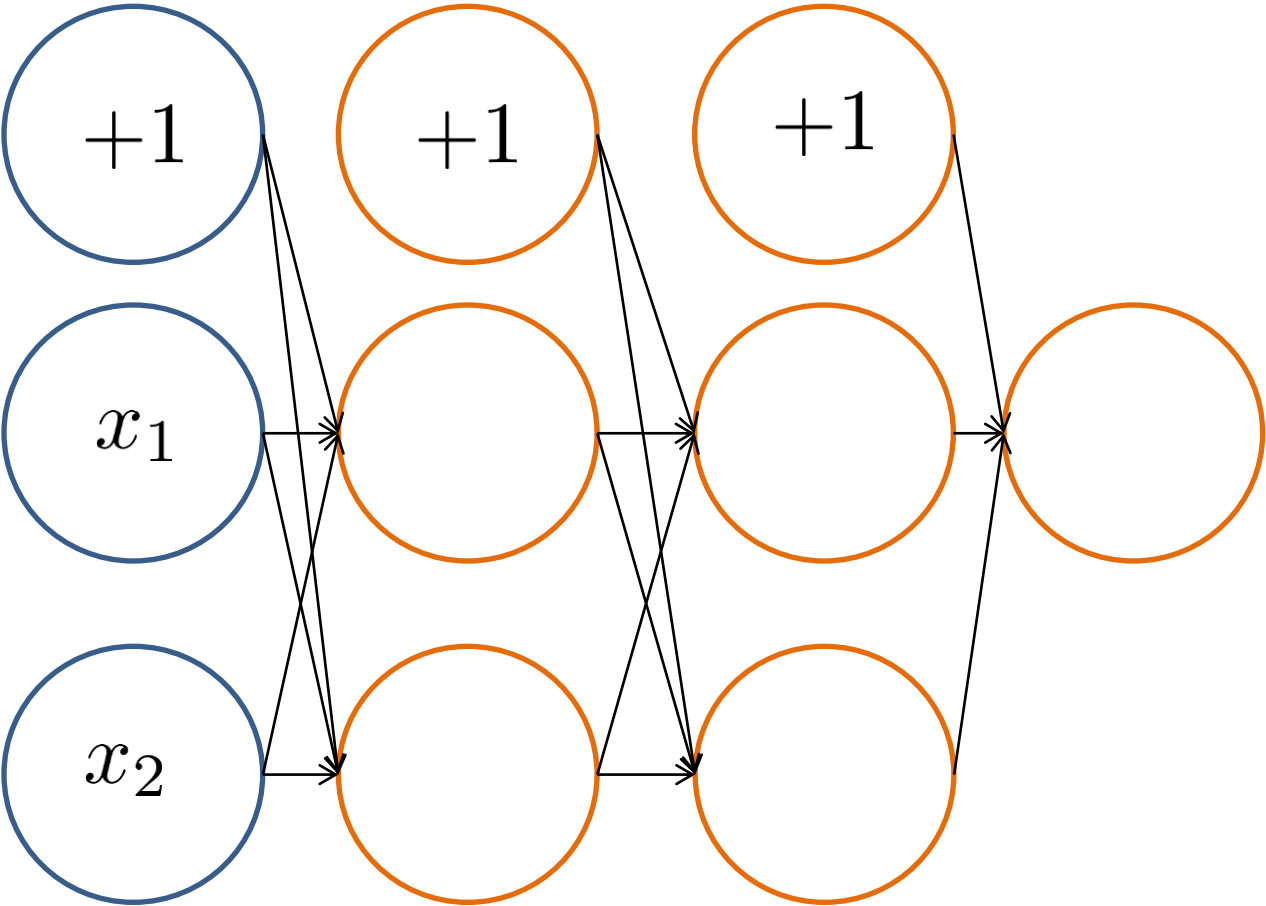
$$D_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)} + \lambda \Theta_{ij}^{(l)} \text{ if } j \neq 0$$

$$D_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)} \text{ if } j = 0$$

$$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = D_{ij}^{(l)}$$

Derivative

# Forward Propagation



# What is backpropagation doing?

$$J(\Theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log(h_{\Theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - (h_{\Theta}(x^{(i)}))) \right] \\ + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (\Theta_{ji}^{(l)})^2$$

Focusing on a single example  $x^{(i)}$ ,  $y^{(i)}$ , the case of 1 output unit, and ignoring regularization ( $\lambda = 0$ ),

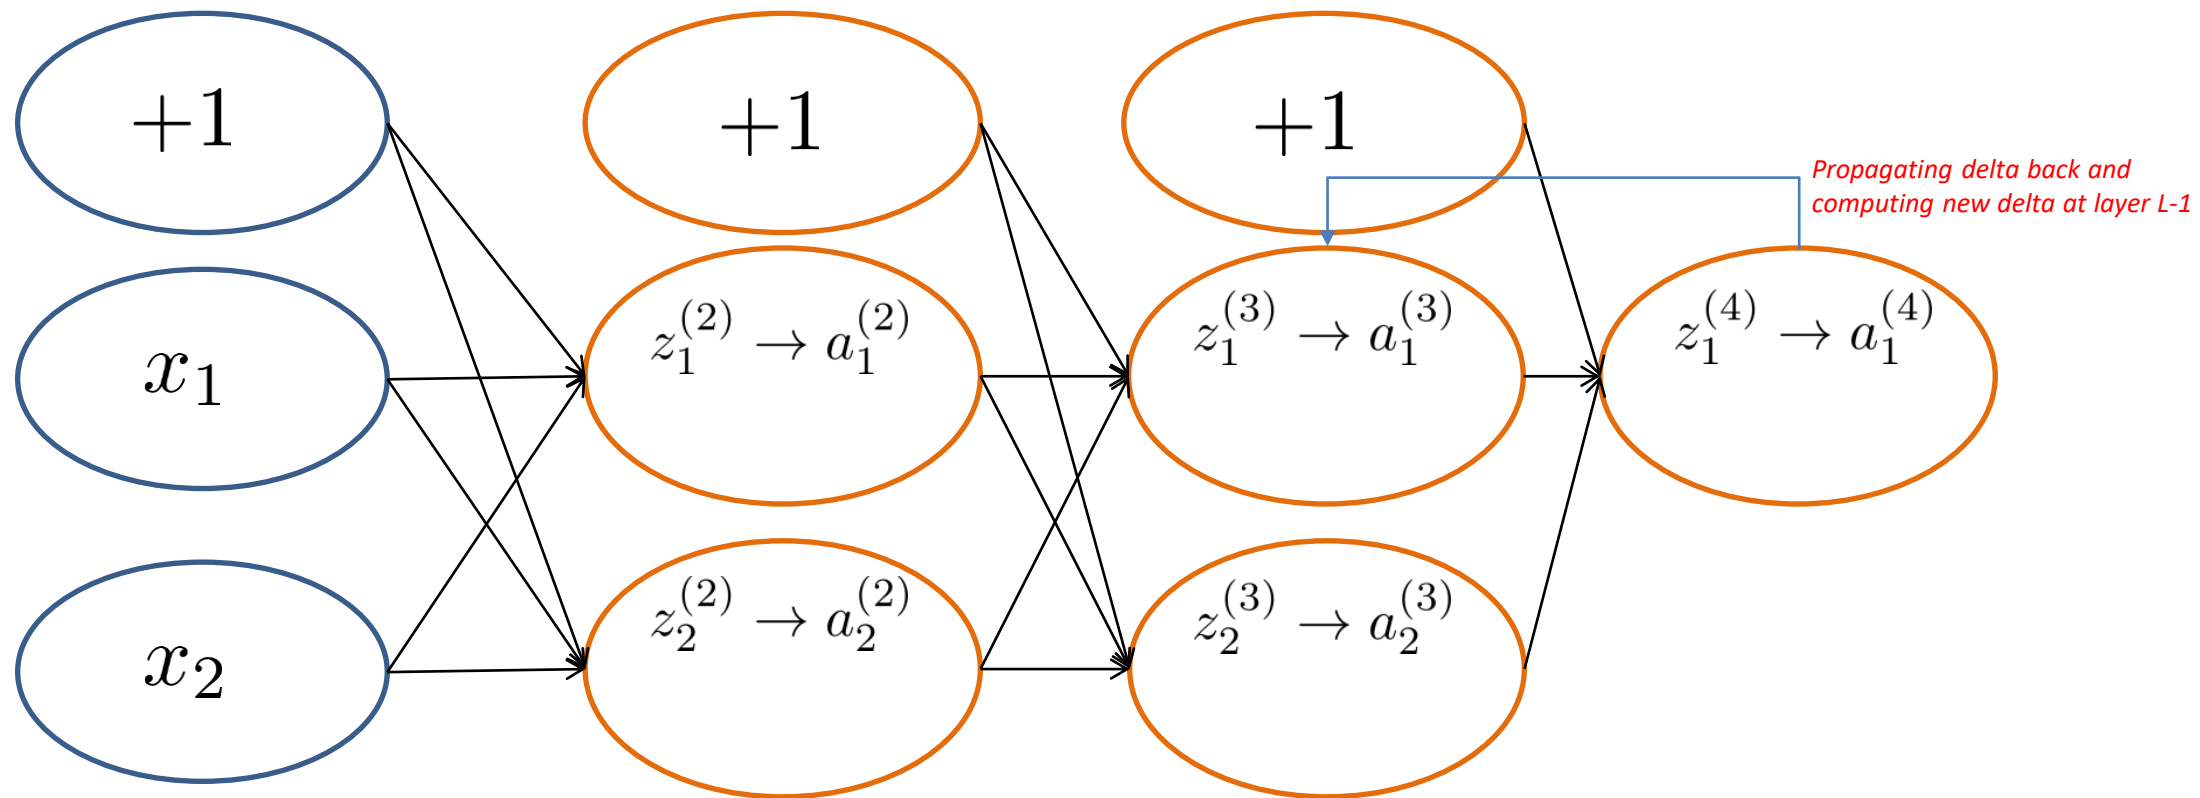
*You can think of cost function as a mean square error function to get a better intuition of back propagation algorithm*

$$\text{cost}(i) = y^{(i)} \log h_{\Theta}(x^{(i)}) + (1 - y^{(i)}) \log h_{\Theta}(x^{(i)})$$

(Think of  $\text{cost}(i) \approx (h_{\Theta}(x^{(i)}) - y^{(i)})^2$ )

I.e. how well is the network doing on example i?

# Forward Propagation



$\delta_j^{(l)}$  = “error” of cost for  $a_j^{(l)}$  (unit  $j$  in layer  $l$ ).

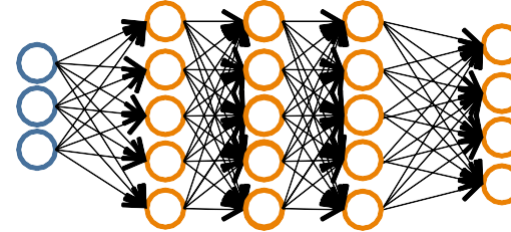
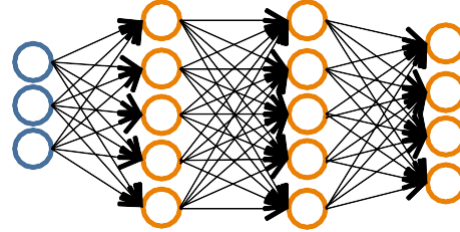
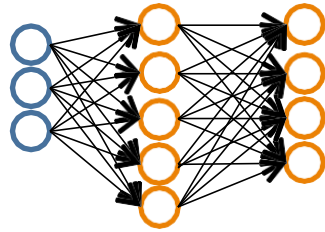
Formally,  $\delta_j^{(l)} = \frac{\partial}{\partial z_j^{(l)}} \text{cost}(i)$  (for  $j \geq 0$ ), where  
 $\text{cost}(i) = y^{(i)} \log h_{\Theta}(x^{(i)}) + (1 - y^{(i)}) \log h_{\Theta}(x^{(i)})$

# Initialization

- For bias
  - Initialize all to 0
- For weights
  - Can't initialize all weights to the same value
    - we can show that all hidden units in a layer will always behave the same
    - need to break symmetry
  - Recipe:  $U[-b, b]$ 
    - the idea is to sample around 0 but break symmetry

# Putting it together

Pick a network architecture



- No. of input units: Dimension of features
- No. output units: Number of classes
- Reasonable default: 1 hidden layer, or if  $>1$  hidden layer, have same no. of hidden units in every layer (usually the more the better)
- Grid search

# Putting it together

## Early stopping

- Use a validation set performance to select the best configuration
- To select the number of epochs, stop training when validation set error increases



# Other tricks of the trade

- Normalizing your (real-valued) data
- Decaying the learning rate
  - as we get closer to the optimum, makes sense to take smaller update steps
- mini-batch
  - can give a more accurate estimate of the risk gradient
- Momentum
  - can use an exponential average of previous gradients



# Dropout

- Idea: «cripple» neural network by removing hidden units
  - each hidden unit is set to 0 with probability 0.5
  - hidden units cannot co-adapt to other units
  - hidden units must be more generally useful